

```
In [64]: import pandas as pd

import numpy as np

import matplotlib.pyplot as plt
%matplotlib inline

import warnings
warnings.filterwarnings("ignore")

import seaborn as sns
```

```
In [65]: data = pd.read_csv('Health Care Diabetes.csv')
```

```
In [66]: data.head()
```

```
Out[66]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

```
In [67]: data.isnull().sum()
```

```
Out[67]: Pregnancies      0
Glucose      0
BloodPressure  0
SkinThickness  0
Insulin      0
BMI          0
DiabetesPedigreeFunction  0
Age          0
Outcome      0
dtype: int64
```

```
In [68]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
Pregnancies      768 non-null int64
Glucose          768 non-null int64
BloodPressure    768 non-null int64
SkinThickness    768 non-null int64
Insulin          768 non-null int64
BMI              768 non-null float64
DiabetesPedigreeFunction  768 non-null float64
Age              768 non-null int64
Outcome          768 non-null int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
In [69]: Positive = data[data['Outcome']==1]
Positive.head(5)
```

```
Out[69]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
2	8	183	64	0	0	23.3	0.672	32	1
4	0	137	40	35	168	43.1	2.288	33	1
6	3	78	50	32	88	31.0	0.248	26	1
8	2	197	70	45	543	30.5	0.158	53	1

Week 1

Data Exploration:


Glucose Descriptive Analysis

```
In [70]: data['Glucose'].value_counts().head(7)
```

```
Out[70]: 100    17
          99    17
          129   14
          125   14
          111   14
          106   14
           95   13
          Name: Glucose, dtype: int64
```

```
In [71]: plt.hist(data['Glucose'])
```

```
Out[71]: (array([ 5.,  0.,  4., 32., 156., 211., 163., 95., 56., 46.]),
          array([ 0., 19.9, 39.8, 59.7, 79.6, 99.5, 119.4, 139.3, 159.2,
                179.1, 199. ]),
          <a list of 10 Patch objects>)
```




Blood Pressure Descriptive Analysis

```
In [72]: data['BloodPressure'].value_counts().head(7)
```

```
Out[72]: 70    57
          74    52
          68    45
          78    45
          72    44
          64    43
          80    40
          Name: BloodPressure, dtype: int64
```

```
In [73]: plt.hist(data['BloodPressure'])
```

```
Out[73]: (array([ 35.,  1.,  2., 13., 107., 261., 243., 87., 14.,  5.]),
          array([ 0., 12.2, 24.4, 36.6, 48.8, 61., 73.2, 85.4, 97.6,
                109.8, 122. ]),
          <a list of 10 Patch objects>)
```




Skin Thickness Descriptive Analysis

```
In [74]: data['SkinThickness'].value_counts().head(7)
```

```
Out[74]: 0    227
          32    31
          30    27
          27    23
          23    22
          33    20
          18    20
          Name: SkinThickness, dtype: int64
```

```
In [75]: plt.hist(data['SkinThickness'])
```

```
Out[75]: (array([231., 107., 165., 175., 78., 9., 2., 0., 0., 1.]),  
array([ 0., 9.9, 19.8, 29.7, 39.6, 49.5, 59.4, 69.3, 79.2, 89.1, 99. ]),  
<a list of 10 Patch objects>)
```




Insulin Descriptive Analysis

```
In [76]: data['Insulin'].value_counts().head(7)
```

```
Out[76]: 0      374  
105      11  
140       9  
130       9  
120       8  
100       7  
94        7  
Name: Insulin, dtype: int64
```

```
In [77]: plt.hist(data['Insulin'])
```

```
Out[77]: (array([487., 155., 70., 30., 8., 9., 5., 1., 2., 1.]),  
array([ 0., 84.6, 169.2, 253.8, 338.4, 423. , 507.6, 592.2, 676.8,  
761.4, 846. ]),  
<a list of 10 Patch objects>)
```




BMI Descriptive Analysis

```
In [78]: data['BMI'].value_counts().head(7)
```

```
Out[78]: 32.0      13  
31.6      12  
31.2      12  
0.0       11  
33.3      10  
32.4      10  
32.8       9  
Name: BMI, dtype: int64
```

```
In [79]: plt.hist(data['BMI'])
```

```
Out[79]: (array([ 11.,  0., 15., 156., 268., 224., 78., 12.,  3.,  1.]),
array([ 0.,  6.71, 13.42, 20.13, 26.84, 33.55, 40.26, 46.97, 53.68,
        60.39, 67.1 ]),
<a list of 10 Patch objects>)
```




```
In [80]: data.describe().transpose()
```

```
Out[80]:
```

	count	mean	std	min	25%	50%	75%	max
Pregnancies	768.0	3.845052	3.369578	0.000	1.00000	3.0000	6.00000	17.00
Glucose	768.0	120.894531	31.972618	0.000	99.00000	117.0000	140.25000	199.00
BloodPressure	768.0	69.105469	19.355807	0.000	62.00000	72.0000	80.00000	122.00
SkinThickness	768.0	20.536458	15.952218	0.000	0.00000	23.0000	32.00000	99.00
Insulin	768.0	79.799479	115.244002	0.000	0.00000	30.5000	127.25000	846.00
BMI	768.0	31.992578	7.884160	0.000	27.30000	32.0000	36.60000	67.10
DiabetesPedigreeFunction	768.0	0.471876	0.331329	0.078	0.24375	0.3725	0.62625	2.42
Age	768.0	33.240885	11.760232	21.000	24.00000	29.0000	41.00000	81.00
Outcome	768.0	0.348958	0.476951	0.000	0.00000	0.0000	1.00000	1.00

```
In [81]: plt.hist(Positive['BMI'],histtype='stepfilled',bins=20)
```

```
Out[81]: (array([ 2.,  0.,  0.,  0.,  0.,  0.,  3., 13., 38., 61., 61., 36., 27.,
        14.,  7.,  3.,  1.,  1.,  0.,  1.]),
array([ 0.,  3.355,  6.71 , 10.065, 13.42 , 16.775, 20.13 , 23.485,
        26.84 , 30.195, 33.55 , 36.905, 40.26 , 43.615, 46.97 , 50.325,
        53.68 , 57.035, 60.39 , 63.745, 67.1 ]),
<a list of 1 Patch objects>)
```




```
In [82]: Positive['BMI'].value_counts().head(7)
```

```
Out[82]: 32.9    8
          31.6    7
          33.3    6
          30.5    5
          32.0    5
          31.2    5
          32.4    4
          Name: BMI, dtype: int64
```

```
In [83]: plt.hist(Positive['Glucose'],histtype='stepfilled',bins=20)
```

```
Out[83]: (array([ 2.,  0.,  0.,  0.,  0.,  0.,  0.,  1.,  4.,  9., 28., 26., 36.,
          27., 29., 22., 24., 21., 25., 14.]),
          array([  0. ,  9.95, 19.9 , 29.85, 39.8 , 49.75, 59.7 , 69.65,
          79.6 , 89.55, 99.5 , 109.45, 119.4 , 129.35, 139.3 , 149.25,
          159.2 , 169.15, 179.1 , 189.05, 199.  ]),
          <a list of 1 Patch objects>)
```




```
In [84]: Positive['Glucose'].value_counts().head(7)
```

```
Out[84]: 125    7
          158    6
          128    6
          115    6
          129    6
          146    5
          162    5
          Name: Glucose, dtype: int64
```

```
In [85]: plt.hist(Positive['BloodPressure'],histtype='stepfilled',bins=20)
```

```
Out[85]: (array([16.,  0.,  0.,  0.,  0.,  1.,  0.,  1.,  6.,  6., 19., 37., 56.,
          36., 41., 31.,  7.,  4.,  4.,  3.]),
          array([  0. ,  5.7, 11.4, 17.1, 22.8, 28.5, 34.2, 39.9, 45.6,
          51.3, 57. , 62.7, 68.4, 74.1, 79.8, 85.5, 91.2, 96.9,
          102.6, 108.3, 114.  ]),
          <a list of 1 Patch objects>)
```




```
In [86]: Positive['BloodPressure'].value_counts().head(7)
```

```
Out[86]: 70    23
          76    18
          78    17
          74    17
          72    16
           0    16
          82    13
          Name: BloodPressure, dtype: int64
```

```
In [87]: plt.hist(Positive['SkinThickness'],histtype='stepfilled',bins=20)
```

```
Out[87]: (array([88.,  1.,  4., 10., 18., 30., 41., 34., 23., 15.,  1.,  1.,  1.,
        0.,  0.,  0.,  0.,  0.,  0.,  1.]),
array([ 0. ,  4.95,  9.9 , 14.85, 19.8 , 24.75, 29.7 , 34.65, 39.6 ,
       44.55, 49.5 , 54.45, 59.4 , 64.35, 69.3 , 74.25, 79.2 , 84.15,
       89.1 , 94.05, 99.  ]),
<a list of 1 Patch objects>)
```




```
In [88]: Positive['SkinThickness'].value_counts().head(7)
```

```
Out[88]: 0      88
        32     14
        33      9
        30      9
        39      8
        35      8
        36      8
        Name: SkinThickness, dtype: int64
```

```
In [89]: plt.hist(Positive['Insulin'],histtype='stepfilled',bins=20)
```

```
Out[89]: (array([141.,  6., 23., 33., 24., 12.,  7.,  7.,  2.,  1.,  1.,
        5.,  3.,  1.,  1.,  0.,  0.,  0.,  0.,  1.]),
array([ 0. , 42.3 , 84.6 , 126.9, 169.2, 211.5, 253.8, 296.1, 338.4,
       380.7, 423. , 465.3, 507.6, 549.9, 592.2, 634.5, 676.8, 719.1,
       761.4, 803.7, 846. ]),
<a list of 1 Patch objects>)
```



```
In [90]: Positive['Insulin'].value_counts().head(7)
```

```
Out[90]: 0      138
        130      6
        180      4
        156      3
        175      3
        194      2
        125      2
        Name: Insulin, dtype: int64
```


Week 2

Data Exploration:


```
In [91]: BloodPressure = Positive['BloodPressure']
Glucose = Positive['Glucose']
SkinThickness = Positive['SkinThickness']
Insulin = Positive['Insulin']
BMI = Positive['BMI']
```

Scatter Plot


```
In [92]: plt.scatter(BloodPressure, Glucose, color='b')
plt.xlabel('BloodPressure')
plt.ylabel('Glucose')
plt.title('BloodPressure & Glucose')
plt.show()
```




```
In [93]: g = sns.scatterplot(x= "Glucose" ,y= "BloodPressure",
hue="Outcome",
data=data);
```



```
In [94]: B = sns.scatterplot(x= "BMI" ,y= "Insulin",
hue="Outcome",
data=data);
```



```
In [95]: S =sns.scatterplot(x= "SkinThickness" ,y= "Insulin",
hue="Outcome",
data=data);
```



Correlation Matrix


```
In [96]: data.corr()
```

```
Out[96]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age
Pregnancies	1.000000	0.129459	0.141282	-0.081672	-0.073535	0.017683	-0.033523	0.544341
Glucose	0.129459	1.000000	0.152590	0.057328	0.331357	0.221071	0.137337	0.263514
BloodPressure	0.141282	0.152590	1.000000	0.207371	0.088933	0.281805	0.041265	0.239528
SkinThickness	-0.081672	0.057328	0.207371	1.000000	0.436783	0.392573	0.183928	-0.113970
Insulin	-0.073535	0.331357	0.088933	0.436783	1.000000	0.197859	0.185071	-0.042163
BMI	0.017683	0.221071	0.281805	0.392573	0.197859	1.000000	0.140647	0.036242
DiabetesPedigreeFunction	-0.033523	0.137337	0.041265	0.183928	0.185071	0.140647	1.000000	0.033561
Age	0.544341	0.263514	0.239528	-0.113970	-0.042163	0.036242	0.033561	1.000000
Outcome	0.221898	0.466581	0.065068	0.074752	0.130548	0.292695	0.173844	0.238356


```
In [97]: sns.heatmap(data.corr())
```

```
Out[97]: <matplotlib.axes._subplots.AxesSubplot at 0x13b945d5f28>
```





```
In [98]: plt.subplots(figsize=(8,8))
sns.heatmap(data.corr(),annot=True,cmap='viridis')
```

Out[98]: <matplotlib.axes._subplots.AxesSubplot at 0x13b9467fc50>



```
In [99]: plt.subplots(figsize=(8,8))
sns.heatmap(data.corr(),annot=True)
```

```
Out[99]: <matplotlib.axes._subplots.AxesSubplot at 0x13b9492ea58>
```



```
In [100]: data.head(5)
```

```
Out[100]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

Week 3

Data Modelling

Logistic Regreation and Model Building

```
In [101]: features = data.iloc[:,[0,1,2,3,4,5,6,7]].values
label = data.iloc[:,8].values
```

Train Test Split

```
In [102]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(features,
label,
test_size=0.2,
random_state =10)
```

Creating Model

```
In [103]: from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
model.fit(X_train,y_train)
```

```
Out[103]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, l1_ratio=None, max_iter=100,
multi_class='auto', n_jobs=None, penalty='l2',
random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
warm_start=False)
```

```
In [104]: print(model.score(X_train,y_train))
print(model.score(X_test,y_test))
```

```
0.7719869706840391
0.7662337662337663
```

```
In [105]: from sklearn.metrics import confusion_matrix
cm = confusion_matrix(label,model.predict(features))
cm
```

```
Out[105]: array([[446,  54],
[122, 146]], dtype=int64)
```

```
In [106]: from sklearn.metrics import classification_report
print(classification_report(label,model.predict(features)))
```


	precision	recall	f1-score	support
0	0.79	0.89	0.84	500
1	0.73	0.54	0.62	268
accuracy			0.77	768
macro avg	0.76	0.72	0.73	768
weighted avg	0.77	0.77	0.76	768

```
In [107]: #Preparing ROC Curve (Receiver Operating Characteristics Curve)
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
```

```
# Predict Probabilities
probs = model.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# Calculating AUC
auc = roc_auc_score(label, probs)
print('AUC: %.3f' % auc)
# Calculating roc curve
fpr, tpr, thresholds = roc_curve(label, probs)
# Plot no skill
plt.plot([0, 1], [0, 1], linestyle='--')
# Plot the roc curve for the model
plt.plot(fpr, tpr, marker='.')
```

```
AUC: 0.837
```

```
Out[107]: [<matplotlib.lines.Line2D at 0x13b94855a58>]
```



Applying Decision Tree Classifier

```
In [108]: from sklearn.tree import DecisionTreeClassifier
model3 = DecisionTreeClassifier(max_depth=5)
model3.fit(X_train,y_train)
```

```
Out[108]: DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
                                max_depth=5, max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, presort='deprecated',
                                random_state=None, splitter='best')
```

```
In [109]: model3.score(X_train,y_train)
```

```
Out[109]: 0.8289902280130294
```

```
In [110]: model3.score(X_test,y_test)
```

```
Out[110]: 0.7727272727272727
```

Applying Random Forest

```
In [111]: from sklearn.ensemble import RandomForestClassifier
model4 = RandomForestClassifier(n_estimators=11)
model4.fit(X_train,y_train)
```

```
Out[111]: RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                                criterion='gini', max_depth=None, max_features='auto',
                                max_leaf_nodes=None, max_samples=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=11,
                                n_jobs=None, oob_score=False, random_state=None,
                                verbose=0, warm_start=False)
```

```
In [112]: model4.score(X_train,y_train)
```

```
Out[112]: 0.988599348534202
```

```
In [113]: model4.score(X_test,y_test)
```

```
Out[113]: 0.7272727272727273
```

Support Vector Classifier

```
In [114]: from sklearn.svm import SVC
model5 = SVC(kernel='rbf', gamma='auto')
model5.fit(X_train,y_train)
```

```
Out[114]: SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
              decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
              max_iter=-1, probability=False, random_state=None, shrinking=True,
              tol=0.001, verbose=False)
```

```
In [115]: model5.score(X_train,y_train)
```

```
Out[115]: 1.0
```

```
In [116]: model5.score(X_test,y_test)
```

```
Out[116]: 0.6168831168831169
```

Applying K-NN

```
In [117]: from sklearn.neighbors import KNeighborsClassifier
model2 = KNeighborsClassifier(n_neighbors=7,
                             metric='minkowski',
                             p = 2)
model2.fit(X_train,y_train)
```

```
Out[117]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                               metric_params=None, n_jobs=None, n_neighbors=7, p=2,
                               weights='uniform')
```

Week 4


Data Modelling:

```
In [118]: #Preparing ROC Curve (Receiver Operating Characteristics Curve)
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score

# Predict Probabilities
probs = model2.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# Calculating AUC
auc = roc_auc_score(label, probs)
print('AUC: %.3f' % auc)
# Calculating roc curve
fpr, tpr, thresholds = roc_curve(label, probs)
print("True Positive Rate - {}, False Positive Rate - {} Thresholds - {}".format(tpr, fpr, thresholds))
# Plot no skill
plt.plot([0, 1], [0, 1], linestyle='--')
# Plot the roc curve for the model
plt.plot(fpr, tpr, marker='.')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
```

```
AUC: 0.836
True Positive Rate - [0.      0.06716418 0.23880597 0.44776119 0.60074627 0.75373134
 0.88059701 0.98507463 1.      ], False Positive Rate - [0.      0.      0.02  0.056 0.12  0.248 0.428 0.668 1.
] Thresholds - [2.      1.      0.85714286 0.71428571 0.57142857 0.42857143
 0.28571429 0.14285714 0.      ]
```

Out[118]: Text(0, 0.5, 'True Positive Rate')




```
In [119]: #Precision Recall Curve for Logistic Regression

from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# Predict Probabilities
probs = model.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# Predict class values
yhat = model.predict(features)
# Calculating precision-recall curve
precision, recall, thresholds = precision_recall_curve(label, probs)
# Calculating F1 score
f1 = f1_score(label, yhat)
# Calculating precision-recall AUC
auc = auc(recall, precision)
# Calculating average precision score
ap = average_precision_score(label, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# Plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# Plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')
```

f1=0.624 auc=0.726 ap=0.727

Out[119]: [<matplotlib.lines.Line2D at 0x13b94b575f8>]




```
In [120]: #Precision Recall Curve for KNN

from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# Predict Probabilities
probs = model2.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# Predict class values
yhat = model2.predict(features)
# Calculating precision-recall curve
precision, recall, thresholds = precision_recall_curve(label, probs)
# Calculating F1 score
f1 = f1_score(label, yhat)
# Calculating precision-recall AUC
auc = auc(recall, precision)
# Calculating average precision score
ap = average_precision_score(label, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# Plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# Plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')
```

f1=0.658 auc=0.752 ap=0.709

Out[120]: [<matplotlib.lines.Line2D at 0x13b94bb73c8>]



```


In [121]: #Precision Recall Curve for Decision Tree Classifier

from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# Predict probabilities
probs = model3.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# Predict class values
yhat = model3.predict(features)
# Calculating precision-recall curve
precision, recall, thresholds = precision_recall_curve(label, probs)
# Calculating F1 score
f1 = f1_score(label, yhat)
# Calculating precision-recall AUC
auc = auc(recall, precision)
# Calculating average precision score
ap = average_precision_score(label, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# Plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# Plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')

```

f1=0.714 auc=0.815 ap=0.768

Out[121]: [<matplotlib.lines.Line2D at 0x13b94be8dd8>]




```
In [122]: #Precision Recall Curve for Random Forest

from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# Predict probabilities
probs = model4.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# Predict class values
yhat = model4.predict(features)
# Calculating precision-recall curve
precision, recall, thresholds = precision_recall_curve(label, probs)
# Calculating F1 score
f1 = f1_score(label, yhat)
# Calculating precision-recall AUC
auc = auc(recall, precision)
# Calculating average precision score
ap = average_precision_score(label, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# Plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# Plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')
```

f1=0.906 auc=0.959 ap=0.950

Out[122]: [<matplotlib.lines.Line2D at 0x13b95dc18d0>]

